

# optimization

July 16, 2025

## 1 Optimization

### 1.1 Imports

```
[344]: import numpy as np
import matplotlib.pyplot as plt
from pymoo.algorithms.moo.nsga2 import NSGA2
from pymoo.algorithms.soo.nonconvex.ga import GA
from pymoo.optimize import minimize
from pymoo.core.problem import Problem
```

### 1.2 Define Model Curves

#### 1.2.1 Pruning vs Accuracy

```
[345]: def get_accuracy(p):
    #6th order polynomial coefficients
    coeffs = [85.28, -10.35, -28.44, -49.40]

    acc = sum(c * p**i for i, c in enumerate(coeffs))

    return max(acc, 0)
```

#### 1.2.2 Pruning vs Model Size

```
[346]: def get_size(p):
    coeffs = [496.7, -706.2, 276.9, 4.020]

    return sum(c * p**i for i, c in enumerate(coeffs))
```

#### 1.2.3 Pruning vs Inference Time

```
[347]: def get_time(p):
    b0 = 0.06386
    b1 = -0.06386

    return (b0 + b1 * p)*10
```

### 1.2.4 Plot Curves

```
[348]: p = np.linspace(0, 1, 100)
sizes = get_size(p)
accuracies = [get_accuracy(p) for p in p]
times = get_time(p)


plt.style.use('default')
plt.figure(figsize=(12, 4))

plt.subplot(1, 3, 1)
plt.plot(p, accuracies, label='Accuracy')
plt.xlabel('Pruning Ratio')
plt.ylabel('Accuracy (%)')
plt.title('Pruning vs Accuracy')
plt.grid()
plt.legend()

plt.subplot(1, 3, 2)
plt.plot(p, sizes, label='Size')
plt.xlabel('Pruning Ratio')
plt.ylabel('Model Size (MB)')
plt.title('Pruning vs Size')
plt.grid()
plt.legend()

plt.subplot(1, 3, 3)
plt.plot(p, times, label='Inference Time')
plt.xlabel('Pruning Ratio')
plt.ylabel('Inference Time (ms)')
plt.title('Pruning vs Inference Time')
plt.grid()
plt.legend()

plt.tight_layout()
plt.show()
```



## 1.3 Rewards

### 1.3.1 Accuracy Rewards

```
[349]: def get_accuracy_reward(curr_accuracy, min_accuracy, sigma_right=4,
    ↪sigma_left=2):
    diff = curr_accuracy - min_accuracy
    if 0<=diff<=1e-2:
        return np.exp(- (diff**2) / (10 * sigma_left**2)) * 100
    else:
        return 1*(np.exp(- (abs(diff)**1.5) / (2 * sigma_right**2)) * 100)
```

### 1.3.2 Inference Time Reward

```
[350]: def get_comp_time_reward(current_comp_time, sigma=0.8):
    return np.exp(- (current_comp_time**2) / (2 * sigma**2))*10
```

### 1.3.3 Model Size Reward

```
[351]: def get_model_size_reward(current_model_size, max_model_size, sigma_left=2):
    diff = current_model_size - max_model_size
    if current_model_size > max_model_size:
        return np.exp(- ((diff)**2) / (10 * sigma_left**2))*99*0.5
    if current_model_size == max_model_size:
        return 99*(0.5)
    else:
        return (99+(current_model_size/max_model_size))*0.5
```

### 1.3.4 Better Pruning Reward

```
[352]: def more_acc_less_size(accuracy, min_accuracy, size, max_model_size):
    if accuracy >= min_accuracy and size <= max_model_size:
        return ((accuracy-min_accuracy)*2) + (max_model_size-size)/2
    return 0
```

### 1.3.5 Final Reward

```
[353]: def get_reward(p, min_accuracy=80.0, max_model_size=350.0, x=10, y=1, z=1) ->
    ↪float:
    accuracy = get_accuracy(p)
    time = get_time(p)
    size = get_size(p)

    acc_reward = np.array(get_accuracy_reward(accuracy, min_accuracy))
    time_reward = np.array(get_comp_time_reward(time))
```

```

size_reward = np.array(get_model_size_reward(size, max_model_size))
better_reward = more_acc_less_size(accuracy, min_accuracy, size,
↪max_model_size)
global counts_of_calulating_rewards
# counts_of_calulating_rewards += 1 # type: ignore

x, y, z = x/(x+y+z), y/(x+y+z), z/(x+y+z)

return (x*acc_reward + y*time_reward + z*size_reward + better_reward + p/2)

```

## 1.4 Importance Values

```

[378]: i_org = np.array([0.03044561, 0.02201545, 0.12376647, 0.09755174, 0.04860051,
0.00832497, 0.03501421, 0.00934147, 0.05674529, 0.15345743,
0.237974 , 0.17676284])
i = i_org.copy()

```

## 1.5 Params

```

[379]: MAX_MODEL_SIZES = [0]*12
GLOBAL_MIN_ACCURACY = 0.0

```

```

[380]: def init():
    global MAX_MODEL_SIZES, GLOBAL_MIN_ACCURACY
    MAX_MODEL_SIZES = [100.0] * 12

    for j in range(12):
        MAX_MODEL_SIZES[j] = MAX_MODEL_SIZES[j] * np.random.uniform(1.0, 5.0)

    GLOBAL_MIN_ACCURACY = 10 * np.random.uniform(5, 8.2)

    print("\n" + "="*60)
    print(f"MAX_MODEL_SIZES: {np.round(MAX_MODEL_SIZES, 2)}")
    print(f"GLOBAL_MIN_ACCURACY: {GLOBAL_MIN_ACCURACY:.2f}%")
    print("="*60)

```

## 1.6 Main Optimization

### 1.6.1 Problem Def

```

[381]: class MultiViewProblem(Problem):
    def __init__(self):
        super().__init__(
            n_var=12,
            n_obj=12,
            xl=np.zeros(12),
            xu=np.ones(12)
        )

```

```

def _evaluate(self, X, out, *args, **kwargs):
    pop_size, _ = X.shape
    F = np.zeros((pop_size, 12))

    for j in range(12):
        ps = X[:, j]
        vals = np.zeros(pop_size)
        for k, p in enumerate(ps):
            # Use global min accuracy for consistency
            r = get_reward(p, GLOBAL_MIN_ACCURACY, MAX_MODEL_SIZES[j])
            weighted_r = r * i[j]
            vals[k] = -weighted_r # minimize negative reward
        F[:, j] = vals

    out["F"] = F

class SingleObjectiveProblem(Problem):
    def __init__(self):
        super().__init__(
            n_var=12,
            n_obj=1,
            xl=np.zeros(12),
            xu=np.ones(12)
        )

    def _evaluate(self, X, out, *args, **kwargs):
        pop_size, _ = X.shape
        F = np.zeros((pop_size, 1))

        for k in range(pop_size):
            p_vec = X[k, :]
            total_reward = 0
            penalty = 0

            weighted_acc = 0
            for j, p in enumerate(p_vec):
                acc = get_accuracy(p)
                size = get_size(p)
                weighted_acc += acc * i[j]

                # Apply penalty for size violations
                if size > MAX_MODEL_SIZES[j]:
                    penalty += (size - MAX_MODEL_SIZES[j]) ** 2 * 100 #

↪ Increased penalty

            r = get_reward(p, GLOBAL_MIN_ACCURACY, MAX_MODEL_SIZES[j])

```

```

        total_reward += r * i[j]

        # Strong penalty if weighted accuracy below global min accuracy
        if weighted_acc < GLOBAL_MIN_ACCURACY:
            penalty += (GLOBAL_MIN_ACCURACY - weighted_acc) ** 2 * 10000 #_
    ↪Much stronger penalty

    F[k, 0] = -(total_reward - penalty)

    out["F"] = F

```

### 1.6.2 Check if feasible

```

[382]: def is_feasible_corrected(p_vec):
    accs = np.array([get_accuracy(p) for p in p_vec])
    sizes = np.array([get_size(p) for p in p_vec])
    weighted_acc = np.sum(accs * i)

    # Check global accuracy constraint
    if weighted_acc < GLOBAL_MIN_ACCURACY:
        return False, accs, sizes

    # Check size constraints
    if np.any(sizes > MAX_MODEL_SIZES):
        return False, accs, sizes

    return True, accs, sizes

```

### 1.6.3 Optimize with relaxed constraints

```

[383]: def optimize_with_relaxed_constraints(min_acc, max_sizes):
    # sourcery skip: low-code-quality

    class RelaxedSingleObjectiveProblem(Problem):
        def __init__(self):
            super().__init__(
                n_var=12,
                n_obj=1,
                xl=np.zeros(12),
                xu=np.ones(12)
            )

        def _evaluate(self, X, out, *args, **kwargs):
            pop_size, _ = X.shape
            F = np.zeros((pop_size, 1))

            for k in range(pop_size):

```

```

        p_vec = X[k, :]
        total_reward = 0
        penalty = 0

        weighted_acc = 0
        for j, p in enumerate(p_vec):
            acc = get_accuracy(p)
            size = get_size(p)
            weighted_acc += acc * i[j]

            # Apply penalty for size violations
            if size > max_sizes[j]:
                penalty += (size - max_sizes[j]) ** 2 * 100

            r = get_reward(p, min_acc, max_sizes[j])
            total_reward += r * i[j]

        # Penalty if weighted accuracy below relaxed min accuracy
        if weighted_acc < min_acc:
            penalty += (min_acc - weighted_acc) ** 2 * 5000

        F[k, 0] = -(total_reward - penalty)

    out["F"] = F

# Run optimization with relaxed constraints
    problem = RelaxedSingleObjectiveProblem()
    algorithm = GA(pop_size=200)

    res = minimize(
        problem,
        algorithm,
        termination=('n_gen', 500),
        verbose=False
    )

    best_p_vec = res.X
    accs = np.array([get_accuracy(p) for p in best_p_vec]) # type: ignore
    sizes = np.array([get_size(p) for p in best_p_vec]) # type: ignore
    weighted_acc = np.sum(accs * i)

    # Check against relaxed constraints
    acc_feasible = weighted_acc >= min_acc
    size_feasible = np.all(sizes <= max_sizes)

    print("\n RELAXED OPTIMIZATION RESULT:")
    print("p values:          ", np.round(best_p_vec, 3)) # type: ignore

```

```

print("view accuracies %:", np.round(accs, 2))
print("weighted acc %:   ", np.round(weighted_acc, 2))
print("model sizes:      ", np.round(sizes, 1))
print("relaxed max sizes:", np.round(max_sizes, 1))
print("original max sizes:", np.round(MAX_MODEL_SIZES, 1))

if acc_feasible and size_feasible:
    print("FEASIBLE with relaxed constraints!")
    total_r = sum(get_reward(p, min_acc, max_sizes[j]) * i[j] for j, p in
↪enumerate(best_p_vec)) # type: ignore
    print("total reward:    ", np.round(total_r, 2))
    return best_p_vec, accs, sizes, weighted_acc, total_r
else:
    print("Still infeasible even with relaxed constraints")

    # Show violations
    if not acc_feasible:
        print(f"Accuracy violation: {weighted_acc:.2f}% < {min_acc:.2f}%")

    if not size_feasible:
        violations = [(j, size, max_sizes[j]) for j, size in
↪enumerate(sizes) if size > max_sizes[j]]
        print(f"Size violations: {len(violations)}")
        for j, size, max_size in violations[:5]: # Show first 5
            print(f"  View {j+1}: {size:.1f} > {max_size:.1f}")

    return best_p_vec, accs, sizes, weighted_acc, None

```

#### 1.6.4 Relax Constraints

```

[384]: def relax_constraints_progressively():
    """Progressively relax constraints to find feasible solutions"""
    print("\n" + "="*60)
    print("PROGRESSIVE CONSTRAINT RELAXATION")
    print("="*60)

    relaxation_scenarios = [
        {"acc_relax": 0, "size_relax": 1.0, "desc": "Original constraints"},
        {"acc_relax": 2, "size_relax": 1.05, "desc": "Relax accuracy by 2%,
↪size by 5%"},
        {"acc_relax": 5, "size_relax": 1.1, "desc": "Relax accuracy by 5%, size
↪by 10%"},
        {"acc_relax": 8, "size_relax": 1.15, "desc": "Relax accuracy by 8%,
↪size by 15%"},
        {"acc_relax": 10, "size_relax": 1.2, "desc": "Relax accuracy by 10%,
↪size by 20%"},
    ]

```



```

for scenario in relaxation_scenarios:
    print(f"\n--- {scenario['desc']} ---")

    # Temporarily adjust constraints
    relaxed_min_acc = GLOBAL_MIN_ACCURACY - scenario['acc_relax']
    relaxed_max_sizes = [size * scenario['size_relax'] for size in
↪MAX_MODEL_SIZES]

    print(f"Target accuracy: {relaxed_min_acc:.1f}%")
    print(f"Max sizes: {np.round(relaxed_max_sizes, 1)}")

    # Test feasibility with relaxed constraints
    feasible_count = 0
    best_solution = None
    best_weighted_acc = 0

    for _ in range(500):
        p_vec = np.random.beta(2, 8, 12)
        accs = np.array([get_accuracy(p) for p in p_vec])
        sizes = np.array([get_size(p) for p in p_vec])
        weighted_acc = np.sum(accs * i)

        # Check relaxed constraints
        acc_feasible = weighted_acc >= relaxed_min_acc
        size_feasible = np.all(sizes <= relaxed_max_sizes)

        if acc_feasible and size_feasible:
            feasible_count += 1
            if weighted_acc > best_weighted_acc:
                best_weighted_acc = weighted_acc
                best_solution = (p_vec, accs, sizes, weighted_acc)

    print(f"Feasible solutions found: {feasible_count}/500 ({feasible_count/
↪5:.1f}%)")

    if feasible_count > 0:
        print(f"Best weighted accuracy found: {best_weighted_acc:.2f}%")

        # If we found feasible solutions, run optimization with relaxed
↪constraints
        if feasible_count >= 10: # If reasonably feasible
            print("Running optimization with relaxed constraints...")
            return optimize_with_relaxed_constraints(relaxed_min_acc,
↪relaxed_max_sizes)
        else:
            print("No feasible solutions found with these constraints.")

```

```

print("\nEven with maximum relaxation, no feasible solutions found.")
print("Constraints may be fundamentally incompatible.")
return None

```

### 1.6.5 Multi Objective Pruning

```

[385]: def optimize_pruning(): # sourcery skip: low-code-quality
    print("Running multi-objective optimization (NSGA-II)...")
    problem = MultiViewProblem()
    algorithm = NSGA2(pop_size=500) # Increased population size

    res = minimize(
        problem,
        algorithm,
        termination=('n_gen', 1000), # Increased generations
        verbose=False
    )

    pareto_ps = res.X

    feasible = []
    for p_vec in pareto_ps: # type: ignore
        valid, accs, sizes = is_feasible_corrected(p_vec)
        if valid:
            feasible.append((p_vec, accs, sizes))

    if feasible:
        best = None
        best_reward = -np.inf
        for p_vec, accs, sizes in feasible:
            total_r = sum(get_reward(p, GLOBAL_MIN_ACCURACY,
↪MAX_MODEL_SIZES[j]) * i[j]
                           for j, p in enumerate(p_vec))
            if total_r > best_reward:
                best_reward = total_r
                best = (p_vec, accs, sizes, total_r)

    p_vec, accs, sizes, total_r = best
    weighted_acc = np.sum(accs * i)

    print("\n FEASIBLE solution found with NSGA-II:")
    print("i values:      ", np.round(i, 3))
    print("p values:      ", np.round(p_vec, 3))
    print("view accuracies %:", np.round(accs, 2))
    print("weighted acc %:  ", np.round(weighted_acc, 2))
    print("model sizes:     ", np.round(sizes, 1))

```

```

print("max sizes:          ", np.round(MAX_MODEL_SIZES, 1))
print("total reward:       ", np.round(total_r, 2))

return p_vec, accs, sizes, weighted_acc, total_r

else:
    print("\nNo feasible solution found with NSGA-II. Trying_
↪single-objective optimization...")
    single_problem = SingleObjectiveProblem()
    single_algorithm = GA(pop_size=200) # Increased population size

    single_res = minimize(
        single_problem,
        single_algorithm,
        termination=('n_gen', 1000), # Increased generations
        verbose=False
    )

    best_p_vec = single_res.X
    valid, accs, sizes = is_feasible_corrected(best_p_vec)

    accs = np.array([get_accuracy(p) for p in best_p_vec])
    sizes = np.array([get_size(p) for p in best_p_vec])
    weighted_acc = np.sum(accs * i)

    if valid:
        total_r = sum(get_reward(p, GLOBAL_MIN_ACCURACY,
↪MAX_MODEL_SIZES[j]) * i[j]
                        for j, p in enumerate(best_p_vec))

        print("\n FEASIBLE solution found with single-objective GA:")
        print("p values:          ", np.round(best_p_vec, 3))
        print("view accuracies %:", np.round(accs, 2))
        print("weighted acc %:      ", np.round(weighted_acc, 2))
        print("model sizes:          ", np.round(sizes, 1))
        print("max sizes:           ", np.round(MAX_MODEL_SIZES, 1))
        print("total reward:        ", np.round(total_r, 2))
        return best_p_vec, accs, sizes, weighted_acc, total_r
    else:
        print("\n Best compromise solution (may violate some constraints):
↪")

        print("p values:          ", np.round(best_p_vec, 3))
        print("view accuracies %:", np.round(accs, 2))
        print("weighted acc %:      ", np.round(weighted_acc, 2))
        print("model sizes:          ", np.round(sizes, 1))
        print("max sizes:           ", np.round(MAX_MODEL_SIZES, 1))

```

```

        # Check constraint violations
        if weighted_acc < GLOBAL_MIN_ACCURACY:
            print(f"\nGlobal accuracy violation: {weighted_acc:.2f}% <
↳{GLOBAL_MIN_ACCURACY}%")

        if size_violations := [
            (j, size, MAX_MODEL_SIZES[j])
            for j, size in enumerate(sizes)
            if size > MAX_MODEL_SIZES[j]
        ]:
            print("\nSize violations:")
            for j, size, max_size in size_violations:
                print(f"  View {j+1}: {size:.1f} > {max_size}")

        print("\n" + "="*60)
        print("STARTING PROGRESSIVE CONSTRAINT RELAXATION...")
        print("="*60)

        # Try progressive relaxation
        relaxation_result = relax_constraints_progressively()
        if relaxation_result is not None:
            return relaxation_result

    return best_p_vec, accs, sizes, weighted_acc, None

```

### 1.6.6 Find a Feasible Starting Point

```

[386]: def find_feasible_starting_point():
        """Find a feasible starting point by working backwards from constraints"""
        print("Searching for feasible starting point...")

        # Start with p values that give high accuracy but might violate size
        ↳constraints
        best_p = None
        best_weighted_acc = 0

        for _ in range(1000):
            # Generate random p values, biased towards lower pruning (higher
            ↳accuracy)
            p_vec = np.random.beta(2, 8, 12) # Biased towards lower values

            accs = np.array([get_accuracy(p) for p in p_vec])
            sizes = np.array([get_size(p) for p in p_vec])
            weighted_acc = np.sum(accs * i)

            # Check if this meets the global accuracy constraint
            if weighted_acc >= GLOBAL_MIN_ACCURACY:

```

```

        # Check size constraints
        if np.all(sizes <= MAX_MODEL_SIZES):
            print(f"Found feasible starting point! Weighted accuracy:␣
↪{weighted_acc:.2f}%")
            return p_vec, accs, sizes, weighted_acc
        else:
            # Track best accuracy even if size constraints are violated
            if weighted_acc > best_weighted_acc:
                best_weighted_acc = weighted_acc
                best_p = p_vec

    if best_p is not None:
        accs = np.array([get_accuracy(p) for p in best_p])
        sizes = np.array([get_size(p) for p in best_p])
        print(f"Best found: Weighted accuracy: {best_weighted_acc:.2f}% (meets␣
↪global min)")
        return best_p, accs, sizes, best_weighted_acc
    else:
        print("No feasible starting point found - constraints may be too tight")
        return None, None, None, None

```

### 1.6.7 Redistribute Accuracy Loss Between Views

```

[387]: def redistribute_accuracy_loss(p_vec, violating_views, target_sizes, i,␣
↪GLOBAL_MIN_ACCURACY):
    # sourcery skip: low-code-quality
    new_p_vec = p_vec.copy()

    # Calculate current weighted accuracy
    current_accs = np.array([get_accuracy(p) for p in p_vec])
    _ = np.sum(current_accs * i)

    # Calculate required pruning for violating views to meet size constraints
    total_accuracy_loss = 0

    for view_idx in violating_views:
        target_size = target_sizes[view_idx]

        # Find pruning amount that gives the target size
        # Use binary search to find p that gives size <= target_size
        p_low, p_high = 0.0, 1.0
        best_p = p_vec[view_idx]

        for _ in range(50): # Binary search iterations
            p_mid = (p_low + p_high) / 2
            size_mid = get_size(p_mid)

```

```

        if size_mid <= target_size:
            best_p = p_mid
            p_high = p_mid
        else:
            p_low = p_mid

        if abs(p_high - p_low) < 1e-6:
            break

        # Calculate accuracy loss for this view
        old_acc = get_accuracy(p_vec[view_idx])
        new_acc = get_accuracy(best_p)
        acc_loss = old_acc - new_acc

        # Weight the accuracy loss by importance
        weighted_acc_loss = acc_loss * i[view_idx]
        total_accuracy_loss += weighted_acc_loss

        # Update pruning for this view
        new_p_vec[view_idx] = best_p

        print(f"View {view_idx+1}: Adjusted pruning from {p_vec[view_idx]:.4f} ↪
        to {best_p:.4f}")
        print(f"    Size: {get_size(p_vec[view_idx]):.1f} ↪ {get_size(best_p):.
        ↪1f} (max: {target_size:.1f})")
        print(f"    Accuracy: {old_acc:.2f}% ↪ {new_acc:.2f}% (loss: {acc_loss:.
        ↪2f}%)")
        print(f"    Weighted loss: {weighted_acc_loss:.2f}%")

    print(f"\nTotal weighted accuracy loss: {total_accuracy_loss:.2f}%")

    # Redistribute the accuracy loss among non-violating views
    non_violating_views = [i for i in range(12) if i not in violating_views]

    if not non_violating_views:
        print("All views violate size constraints - cannot redistribute")
        return new_p_vec, False

    # Calculate how much we need to recover in weighted accuracy
    new_weighted_acc = np.sum([get_accuracy(new_p_vec[j]) * i[j] for j in ↪
    ↪range(12)])
    accuracy_deficit = GLOBAL_MIN_ACCURACY - new_weighted_acc

    print(f"New weighted accuracy after size adjustments: {new_weighted_acc:.
    ↪2f}%")
    print(f"Required global minimum: {GLOBAL_MIN_ACCURACY:.2f}%")
    print(f"Accuracy deficit to recover: {accuracy_deficit:.2f}%")

```

```

if accuracy_deficit <= 0:
    print("No redistribution needed - global constraint already satisfied")
    return new_p_vec, True

# Try to recover the deficit by reducing pruning in non-violating views
print(f"\nRedistributing {accuracy_deficit:.2f}% accuracy deficit among_
↪{len(non_violating_views)} views...")

# Sort non-violating views by their ability to improve (lower current_
↪pruning = more room to improve)
view_flexibility = [(j, new_p_vec[j]) for j in non_violating_views]
view_flexibility.sort(key=lambda x: x[1]) # Sort by current pruning_
↪(ascending)

remaining_deficit = accuracy_deficit

for view_idx, current_p in view_flexibility:
    if remaining_deficit <= 0:
        break

    # Calculate how much this view can contribute
    # Try reducing pruning to improve accuracy
    min_p = 0.0
    max_improvement_p = current_p

    if max_improvement_p >= min_p:
        # Calculate potential accuracy improvement
        current_acc = get_accuracy(current_p)
        improved_acc = get_accuracy(max_improvement_p)
        potential_improvement = (improved_acc - current_acc) * i[view_idx]

        # Take what we need or what's available, whichever is smaller
        if potential_improvement > 0:
            # Calculate exact pruning needed for this view's contribution
            needed_acc_improvement = min(remaining_deficit / i[view_idx],_
↪improved_acc - current_acc)

            # Find pruning that gives this accuracy improvement
            target_acc = current_acc + needed_acc_improvement

            # Binary search for pruning that gives target accuracy
            p_low, p_high = 0, current_p
            best_p = current_p

            for _ in range(50):
                p_mid = (p_low + p_high) / 2

```

```

        acc_mid = get_accuracy(p_mid)

        if acc_mid >= target_acc:
            best_p = p_mid
            p_high = p_mid
        else:
            p_low = p_mid

        if abs(p_high - p_low) < 1e-6:
            break

        # Check if this adjustment violates size constraint
        new_size = get_size(best_p)
        if new_size <= MAX_MODEL_SIZES[view_idx]:
            actual_improvement = (get_accuracy(best_p) - current_acc) * i[view_idx]

            new_p_vec[view_idx] = best_p
            remaining_deficit -= actual_improvement

            print(f"View {view_idx+1}: Reduced pruning from {current_p:.4f} to {best_p:.4f}")
            print(f"    Accuracy: {current_acc:.2f}% → {get_accuracy(best_p):.2f}%")
            print(f"    Weighted contribution: {actual_improvement:.2f}%")
            print(f"    Size: {get_size(current_p):.1f} → {new_size:.1f} (max: {MAX_MODEL_SIZES[view_idx]:.1f})")
        else:
            print(f"View {view_idx+1}: Cannot reduce pruning further without violating size constraint")

        # Check final result
        final_weighted_acc = np.sum([get_accuracy(new_p_vec[j]) * i[j] for j in range(12)])
        final_success = final_weighted_acc >= GLOBAL_MIN_ACCURACY

        print(f"\n{'='*50}")
        print("REDISTRIBUTION RESULT:")
        print(f"Final weighted accuracy: {final_weighted_acc:.2f}%")
        print(f"Global minimum required: {GLOBAL_MIN_ACCURACY:.2f}%")
        print(f"Success: {'Y' if final_success else 'N'}")

        if not final_success:
            print(f"Remaining deficit: {GLOBAL_MIN_ACCURACY - final_weighted_acc:.2f}%")

        return new_p_vec, final_success

```



### 1.6.8 Handle Size Violations

```
[388]: def handle_size_violations(p_vec, i, global_min_accuracy, MAX_MODEL_SIZES):
    print("\n" + "="*60)
    print("HANDLING SIZE CONSTRAINT VIOLATIONS")
    print("="*60)

    # Identify violating views
    current_sizes = np.array([get_size(p) for p in p_vec])
    violating_views = [i for i, size in enumerate(current_sizes) if size >
↪MAX_MODEL_SIZES[i]]

    if not violating_views:
        print("No size violations detected.")
        accs = np.array([get_accuracy(p) for p in p_vec])
        weighted_acc = np.sum(accs * i)
        return p_vec, True, {
            'weighted_accuracy': weighted_acc,
            'view accuracies': accs,
            'view sizes': current_sizes,
            'violations': []
        }

    print(f"Found {len(violating_views)} views violating size constraints:")
    for view_idx in violating_views:
        print(f"  View {view_idx+1}: {current_sizes[view_idx]:.1f} >
↪{MAX_MODEL_SIZES[view_idx]:.1f}")

    # Redistribute accuracy loss
    adjusted_p_vec, success = redistribute_accuracy_loss(
        p_vec, violating_views, MAX_MODEL_SIZES, i, global_min_accuracy
    )

    # Calculate final metrics
    final_accs = np.array([get_accuracy(p) for p in adjusted_p_vec])
    final_sizes = np.array([get_size(p) for p in adjusted_p_vec])
    final_weighted_acc = np.sum(final_accs * i)

    remaining_violations = [
        {
            'view': j + 1,
            'size': final_sizes[j],
            'max_size': MAX_MODEL_SIZES[j],
            'violation': final_sizes[j] - MAX_MODEL_SIZES[j],
        }
        for j in range(12)
        if final_sizes[j] > MAX_MODEL_SIZES[j]
```

```

]
final_metrics = {
    'weighted_accuracy': final_weighted_acc,
    'view accuracies': final_accs,
    'view sizes': final_sizes,
    'violations': remaining_violations,
    'global_constraint_met': final_weighted_acc >= global_min_accuracy
}

return adjusted_p_vec, success and not remaining_violations, final_metrics

```

### 1.6.9 Check if redistribution was successful

```

[389]: def successful_redist(adjusted_p_vec, i, metrics, total_r):
    # Calculate new total reward
    new_total_r = sum(get_reward(p, GLOBAL_MIN_ACCURACY, MAX_MODEL_SIZES[j]) *
↪ i[j]
                        for j, p in enumerate(adjusted_p_vec))

    print("\nSUCCESSFUL REDISTRIBUTION:")
    print("Adjusted p values:      ", np.round(adjusted_p_vec, 3))
    print("View accuracies %:       ", np.round(metrics['view accuracies'], 4))
    print("Weighted accuracy %:      ", np.round(metrics['weighted_accuracy'], 4))
    print("Model sizes:              ", np.round(metrics['view sizes'], 2))
    print("Max sizes:                 ", np.round(MAX_MODEL_SIZES, 2))
    print("New total reward:          ", np.round(new_total_r, 2))
    print("Original total reward:", np.round(total_r, 2))

    return adjusted_p_vec, metrics['view accuracies'], metrics['view sizes'],
↪ metrics['weighted_accuracy'], new_total_r

```

### 1.6.10 Optimize with Redist

```

[390]: def optimize_pruning_with_redistribution():
    print("Running enhanced optimization with redistribution...")

    # First try standard optimization
    result = optimize_pruning()

    if result is None:
        print("Standard optimization failed. Trying redistribution approach...")
        return None

    p_vec, accs, sizes, weighted_acc, total_r = result

    # Fix: Use proper variable names to avoid conflicts

```

```

size_violations = [view_idx for view_idx, size in enumerate(sizes) if size_
↪ MAX_MODEL_SIZES[view_idx]] # type: ignore

if not size_violations:
    print("No size violations - solution is already feasible!")
    return result

print(f"\nDetected {len(size_violations)} size violations. Applying_
↪ redistribution...")

# Debug: Print the actual violations
print("Size violations detected:")
for view_idx in size_violations:
    print(f"  View {view_idx+1}: {sizes[view_idx]:.1f} >_
↪ {MAX_MODEL_SIZES[view_idx]:.1f}") # type: ignore

# Apply redistribution
adjusted_p_vec, success, metrics = handle_size_violations(
    p_vec, i, GLOBAL_MIN_ACCURACY, MAX_MODEL_SIZES # type: ignore
)

if success:
    return successful_redist(adjusted_p_vec, i, metrics, total_r)

print("Redistribution failed - returning original solution with violations")
return result

```

### 1.6.11 Params

[409]: `init()`

```

=====
MAX_MODEL_SIZES: [396.94 446.72 236.25 183.97 483.84 398.09 100.33 137.14 318.26
460.51
468.65 294.01]
GLOBAL_MIN_ACCURACY: 64.59%
=====

```

### 1.6.12 Main Loop

[410]: `print("Testing redistribution approach...")`  
`print("="*70)`  
`print("Global minimum accuracy:", GLOBAL_MIN_ACCURACY)`  
`print("Max model sizes:", np.round(MAX_MODEL_SIZES, 2))`  
`print("="*70)`

```

test_result = find_feasible_starting_point()
if test_result[0] is not None:
    print("Running optimization with redistribution...")
    final_result = optimize_pruning_with_redistribution()

    if final_result is not None:
        print("\nFINAL OPTIMIZED SOLUTION WITH REDISTRIBUTION COMPLETE")
        print("Adjusted p values:      ", np.round(final_result[0], 3))
        print("View accuracies %:      ", np.round(final_result[1], 2))
        print("Weighted accuracy %:      ", np.round(final_result[3], 2))
        print("Required global accuracy: ", GLOBAL_MIN_ACCURACY)
        print("Model sizes:      ", np.round(final_result[2], 1))
        print("Max sizes:      ", np.round(MAX_MODEL_SIZES, 1))
        print("Total reward:      ", np.round(final_result[4], 2)) if
↪final_result[4] is not None else print("No total reward calculated.")

        print(f"{'View':<6} {'Size':<10} {'Max Size':<12} {'Accuracy (%)':<15}
↪{'Importance':<10}, {'Prune':<10}")
        print("-" * 70)

        views_violated = []

        for view in range(12):
            size = final_result[2][view]
            max_size = MAX_MODEL_SIZES[view]
            acc = final_result[1][view]
            importance = i[view]
            prune = final_result[0][view]
            if np.round(size, 1) > np.round(max_size, 1):
                views_violated.append(view + 1)
            print(f"{'view+1':<6}, {'size':<10.1f}, {'max_size':<12.1f}, {'acc':<15.
↪2f}, {'importance*100':<12.2f}, {'prune':<10.3f}")
            print("\n" + "="*70)
            print("Final Weighted Accuracy: ", np.round(final_result[3], 1),
↪"Required:", np.round(GLOBAL_MIN_ACCURACY, 1))
            print("Views violating size constraints:", views_violated if
↪views_violated else "None")
            if np.round(final_result[3], 1) >= np.round(GLOBAL_MIN_ACCURACY, 1):
                print("Solution meets global accuracy requirement!")
            else:
                print("Solution does NOT meet global accuracy requirement.")

        else:
            print("Constraints appear too tight even for redistribution approach.")

```

Testing redistribution approach...

=====

Global minimum accuracy: 64.58665441584051  
 Max model sizes: [396.94 446.72 236.25 183.97 483.84 398.09 100.33 137.14 318.26  
 460.51  
 468.65 294.01]

=====

Searching for feasible starting point...  
 Best found: Weighted accuracy: 84.30% (meets global min)  
 Running optimization with redistribution...  
 Running enhanced optimization with redistribution...  
 Running multi-objective optimization (NSGA-II)...

No feasible solution found with NSGA-II. Trying single-objective optimization...

Best compromise solution (may violate some constraints):  
 p values: [0.494 0.501 0.503 0.572 0.507 0.502 0.844 0.707 0.492 0.5  
 0.502 0.499]  
 view accuracies %: [67.28 66.75 66.61 60.78 66.25 66.66 26.58 46.25 67.4 66.8  
 66.64 66.89]  
 weighted acc %: 64.59  
 model sizes: [216. 212.9 212.2 184. 210.2 212.5 100.3 137.1 216.6 213.2  
 212.3 213.7]  
 max sizes: [396.9 446.7 236.2 184. 483.8 398.1 100.3 137.1 318.3 460.5  
 468.7 294. ]

Global accuracy violation: 64.59% < 64.58665441584051%

=====

STARTING PROGRESSIVE CONSTRAINT RELAXATION...

=====

=====

PROGRESSIVE CONSTRAINT RELAXATION

=====

--- Original constraints ---

Target accuracy: 64.6%  
 Max sizes: [396.9 446.7 236.2 184. 483.8 398.1 100.3 137.1 318.3 460.5 468.7  
 294. ]  
 Feasible solutions found: 0/500 (0.0%)  
 No feasible solutions found with these constraints.

--- Relax accuracy by 2%, size by 5% ---

Target accuracy: 62.6%  
 Max sizes: [416.8 469.1 248.1 193.2 508. 418. 105.3 144. 334.2 483.5 492.1  
 308.7]  
 Feasible solutions found: 0/500 (0.0%)  
 No feasible solutions found with these constraints.

--- Relax accuracy by 5%, size by 10% ---

Target accuracy: 59.6%

Max sizes: [436.6 491.4 259.9 202.4 532.2 437.9 110.4 150.9 350.1 506.6 515.5 323.4]

Feasible solutions found: 0/500 (0.0%)

No feasible solutions found with these constraints.

--- Relax accuracy by 8%, size by 15% ---

Target accuracy: 56.6%

Max sizes: [456.5 513.7 271.7 211.6 556.4 457.8 115.4 157.7 366. 529.6 538.9 338.1]

Feasible solutions found: 0/500 (0.0%)

No feasible solutions found with these constraints.

--- Relax accuracy by 10%, size by 20% ---

Target accuracy: 54.6%

Max sizes: [476.3 536.1 283.5 220.8 580.6 477.7 120.4 164.6 381.9 552.6 562.4 352.8]

Feasible solutions found: 0/500 (0.0%)

No feasible solutions found with these constraints.

Even with maximum relaxation, no feasible solutions found.

Constraints may be fundamentally incompatible.

No size violations - solution is already feasible!

#### FINAL OPTIMIZED SOLUTION WITH REDISTRIBUTION COMPLETE

Adjusted p values: [0.494 0.501 0.503 0.572 0.507 0.502 0.844 0.707 0.492 0.5 0.502 0.499]

View accuracies %: [67.28 66.75 66.61 60.78 66.25 66.66 26.58 46.25 67.4 66.8 66.64 66.89]

Weighted accuracy %: 64.59

Required global accuracy: 64.58665441584051

Model sizes: [216. 212.9 212.2 184. 210.2 212.5 100.3 137.1 216.6 213.2 212.3 213.7]

Max sizes: [396.9 446.7 236.2 184. 483.8 398.1 100.3 137.1 318.3 460.5 468.7 294. ]

No total reward calculated.

View	Size	Max Size	Accuracy (%)	Importance	Prune
1	, 216.0	, 396.9	, 67.28	, 3.04	, 0.494
2	, 212.9	, 446.7	, 66.75	, 2.20	, 0.501
3	, 212.2	, 236.2	, 66.61	, 12.38	, 0.503
4	, 184.0	, 184.0	, 60.78	, 9.76	, 0.572
5	, 210.2	, 483.8	, 66.25	, 4.86	, 0.507
6	, 212.5	, 398.1	, 66.66	, 0.83	, 0.502
7	, 100.3	, 100.3	, 26.58	, 3.50	, 0.844
8	, 137.1	, 137.1	, 46.25	, 0.93	, 0.707
9	, 216.6	, 318.3	, 67.40	, 5.67	, 0.492

10	, 213.2	, 460.5	, 66.80	, 15.35	, 0.500
11	, 212.3	, 468.7	, 66.64	, 23.80	, 0.502
12	, 213.7	, 294.0	, 66.89	, 17.68	, 0.499

```
=====
Final Weighted Accuracy: 64.6 Required: 64.6
Views violating size constraints: None
Solution meets global accuracy requirement!
```

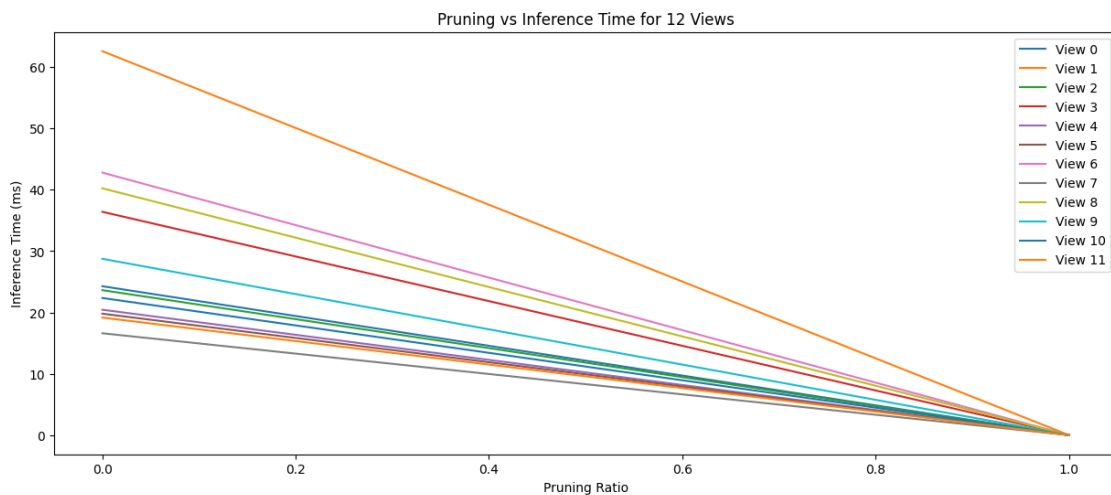
## 1.7 Optimizing Computation Time

```
[368]: times = {}

plt.figure(figsize=(15, 6))
for v in range(12):
    k = np.random.randint(10, 100)
    t = [(get_time(p)*k) for p in np.linspace(0, 1, 100)]
    plt.plot(np.linspace(0, 1, 100), t, label=f'View {v}')
    times[v] = np.array(t)
    # print(f"View {v+1} inference time: {t[-50]:.2f} ms at p=0.5")

plt.xlabel('Pruning Ratio')
plt.ylabel('Inference Time (ms)')
plt.title('Pruning vs Inference Time for 12 Views')
plt.legend()
```

[368]: <matplotlib.legend.Legend at 0x169da9f90>



```
[369]: f_p = np.array([
    0.1, 0.5, 0.3, 0.4, 0.6, 0.2, 0.8, 0.9, 0.1, 0.7, 0.15, 0.32
```

```

])

bottleneck = 0.0
view_bottleneck = 0

for v in range(12):
    t = times[v]
    p = f_p[v]
    new_b = t[int(p*100)]

    if new_b > bottleneck:
        bottleneck = new_b
        view_bottleneck = v

bottleneck, view_bottleneck

```

```
[369]: (np.float64(36.1679818181815), 8)
```

```
[370]: times_at_p = np.array([
        times[v][int(f_p[v]*100)] for v in range(12)
    ])

```

```
[412]: times_at_p.sort()
        times_at_p.mean()

```

```
[412]: np.float64(16.753036700336704)
```

```
[ ]: !jupyter nbconvert --to pdf optimization.ipynb

```